

Bachelor's thesis

# **PROJECT ADAPTATION IN CODE COMPLETION VIA IN-CONTEXT LEARNING**

**Maksim Sapronov**

Faculty of Information Technology  
Department of Applied Mathematics  
Supervisor: Evgenii Glukhov M.Sc.  
May 6, 2025



## Assignment of bachelor's thesis

**Title:** Project adaptation in code completion via in-context learning  
**Student:** Maksim Sapronov  
**Supervisor:** Evgenii Glukhov, M.Sc.  
**Study program:** Informatics  
**Branch / specialization:** Artificial Intelligence 2021  
**Department:** Department of Applied Mathematics  
**Validity:** until the end of summer semester 2025/2026

### Instructions

In recent years, advancements in the fields of code and natural language processing have significantly transformed code completion systems, making them more accurate and capable of assisting with software development tasks. Some recent Large Language Models (LLMs) and Code LLMs can utilize an entire software project. However, most open source pre-trained models can process less than 5% of the project's code files at once. Such a restricted context length of Code LLMs limits their ability to comprehend the structure of a real-world software project fully. As a result, models often struggle to integrate information spread across multiple files, such as dependencies, class hierarchies, and external library usage, leading to outputs that lack project-wide coherence.

It is known that LLMs have an ability for in-context learning, i.e. it can adapt to novel tasks if there are some examples in the context. In particular, for the code completion task, some of the project files are included in the context as examples of completion. It is common practice now to include a repository-level training step in the pre-training of Code LLMs. However, some open questions in this domain are the main focus of this research project.

Research Questions:

For pre-trained Code LLMs with no context window extension:

1. Does the improvement of the quality of code completion depend on a context



composition approach for in-context learning? Estimate the influence.

2. Does fine-tuning help to improve the quality of code completion for a particular context composer?

For pre-trained Code LLMs with context window extension:

1. Are there any context composition techniques that don't impact in-context learning abilities?

2. Does the improvement of the quality of code completion depend on a context composition approach for the repository-level training step? Estimate the influence.

Guidelines for Elaboration:

1. Conduct a survey of the existing papers on the following topics: in-context learning, code completion models, repository-level code completion.

2. Developing a Context Composition Framework for Repository-level Code Completion Task

- Design and implement a framework for extracting a context from a repository, include language-specific processing for code files.

3. Investigating the Impact of Context Composition on Code Completion

- Analyze how variations in context composition (from 2) and context length influence the quality of one-line code completion for a pre-trained Code LLM.

4. Developing a Fine-Tuning Pipeline

- Design and implement a pipeline for the complete fine-tuning of a pre-trained Code LLM.

5. Evaluate Context Composition Methods on Fine-tuned Model

- Base model is <https://huggingface.co/deepseek-ai/deepseek-coder-1.3b-base>

- The dataset for the fine-tuning is available to company staff at JetBrains through internal resources. It contains data from open-source repositories on GitHub.

6. Extend Model's Context Length with Repository-level Training

- Base model is <https://opencoder-llm.github.io>

7. Compare Context Extensions for Different Context Composers

- Dataset for the fine-tuning is available to company staff at JetBrains through internal resources. It contains data from open-source repositories on GitHub.

- For evaluation use the approach that is suggested in <https://arxiv.org/abs/2406.11612>

Czech Technical University in Prague

Faculty of Information Technology

© 2025 Maksim Sapronov. All rights reserved.

*This thesis is school work as defined by Copyright Act of the Czech Republic. It has been submitted at Czech Technical University in Prague, Faculty of Information Technology. The thesis is protected by the Copyright Act and its usage without author's permission is prohibited (with exceptions defined by the Copyright Act).*

Citation of this thesis: Sapronov Maksim. *Project adaptation in code completion via in-context learning*. Bachelor's thesis. Czech Technical University in Prague, Faculty of Information Technology, 2025.

*Thank you . . .*

## Declaration

I hereby declare that the presented thesis is my own work and that I have cited all sources of information in accordance with the Guideline for adhering to ethical principles when elaborating an academic final thesis.

I acknowledge that my thesis is subject to the rights and obligations stipulated by the Act No. 121/2000 Coll., the Copyright Act, as amended. In accordance with Section 2373(2) of Act No. 89/2012 Coll., the Civil Code, as amended, I hereby grant a non-exclusive authorization (license) to utilize this thesis, including all computer programs that are part of it or attached to it and all documentation thereof (hereinafter collectively referred to as the “Work”), to any and all persons who wish to use the Work. Such persons are entitled to use the Work in any manner that does not diminish the value of the Work and for any purpose (including use for profit). This authorization is unlimited in time, territory and quantity.

The use of generative AI tools, such as GPT-4o, Claude 3.7 Sonnet, and DeepL, was strictly limited to checking grammar and rephrasing pre-written text without altering its original meaning. These tools were not utilized as sources of information on the subject of this Work. Furthermore, Perplexity services were employed for search purposes, with subsequent human analysis of the results. I accept full responsibility for the outcomes associated with the use of these tools.

In Prague on May 6, 2025

## Abstract

Fill in the abstract of this thesis in English.

**Keywords** code completion, repository-level code completion, project adaptation, in-context learning, long context, context extension, Transformer, Code LLM

## Abstrakt

Fill in the abstract of this thesis in Czech.

**Klíčová slova** dokončování kódu, dokončování kódu na úrovni repozitáře, projektové přizpůsobení, učení v kontextu, dlouhý kontext, rozšíření kontextu, Transformátor, Code LLM

## Contents

<b>Objectives of the Thesis</b>	<b>1</b>
<b>Introduction</b>	<b>2</b>
<b>I Conceptual Framework</b>	<b>4</b>
<b>1 Code Completion</b>	<b>5</b>
1.1 Task Definition . . . . .	5
1.2 Motivation . . . . .	5
1.3 Evolution of Methods . . . . .	6
1.4 Taxonomy of Code Completion . . . . .	8
<b>2 Standard LM</b>	<b>11</b>
2.1 Definition . . . . .	11
2.2 Text Representation . . . . .	12
2.2.1 Tokens . . . . .	12
2.2.2 Byte Pair Encoding for Tokenization . . . . .	13
2.2.3 Embeddings . . . . .	13
2.3 Transformer Models . . . . .	14
2.3.1 Overview of the Forward Pass . . . . .	15
2.3.2 Decoder Layer . . . . .	15
2.3.3 Normalization . . . . .	16
2.3.4 Attention . . . . .	16
2.3.5 Multi-Head Attention . . . . .	16
2.3.6 Multi-Layer Perceptron . . . . .	17
2.3.7 Dropout . . . . .	17
2.3.8 Rotary Position Embedding . . . . .	17
2.4 Training . . . . .	18
2.4.1 Maximum Likelihood Estimation . . . . .	18
2.4.2 Gradient-Based Optimization . . . . .	19
2.4.3 Batch-Related Terminology . . . . .	19
2.4.4 Learning Rate Scheduling . . . . .	20
2.5 Inference . . . . .	20
2.5.1 Sampling . . . . .	20
2.5.2 Stopping Criteria . . . . .	20



<b>3</b>	<b>Completion-Centric LM</b>	<b>22</b>
3.1	Training Stages . . . . .	22
3.2	Gradient Masking . . . . .	23
3.3	Fill-in-the-Middle . . . . .	23
3.4	Evaluation . . . . .	23
3.4.1	Metrics . . . . .	23
3.4.2	Evaluation Modes . . . . .	27
3.4.3	Benchmarks . . . . .	27
3.5	Modern Code Completion Models . . . . .	30
<b>4</b>	<b>In-Context Learning</b>	<b>31</b>
<b>5</b>	<b>Long Context</b>	<b>32</b>
<b>II</b>	<b>Applied Research</b>	<b>33</b>
<b>6</b>	<b>Technical Foundation</b>	<b>34</b>
6.1	Context Composition Framework . . . . .	34
6.1.1	Building Blocks . . . . .	34
6.1.2	List of Context Composers . . . . .	35
6.1.3	Order Modifications . . . . .	37
6.1.4	Examples . . . . .	37
6.2	Fine-Tuning Pipeline . . . . .	38
6.3	Tools . . . . .	39
<b>7</b>	<b>Research Investigation</b>	<b>40</b>
7.1	Research Questions . . . . .	40
7.2	Experimental Design . . . . .	40
7.2.1	Training Data . . . . .	41
7.2.2	Training . . . . .	41
7.2.3	Evaluation . . . . .	41
7.3	Composition Impact on Inference . . . . .	41
7.4	Fine-Tuning on Compositions . . . . .	41
7.5	Effect of Context Extension . . . . .	41
7.6	Influence of Composition on Context Extension . . . . .	41
7.7	Limitations and Future Work . . . . .	41
	<b>Conclusion</b>	<b>42</b>
<b>A</b>	<b>Appendix</b>	<b>44</b>
	<b>Contents of the attachment</b>	<b>51</b>

List of Figures

6.1 Order of composer blocks. Nodes represent block types: ① File Filtering, ② File Preprocessing, ③ File Chunking, ④ Chunk Ranking, ⑤ Chunk Sorting, ⑥ Chunk Assembling, and ⑦ Context Postprocessing. Edges indicate the permissible sequence of block application. The data format is denoted by three frames: Files, Chunks, and Context. The bold subgraph highlights blocks that may be omitted. . . . . 35

6.2 Three modes of context ordering. The numbers in boxes indicate the rank of the chunks, with smaller numbers denoting higher relevance. . . . . 37

List of Tables

List of code listings

TODO: “Supporting visualizations will be added after the main text is complete.”

## List of abbreviations

Will Be Done at the Very End

# Objectives of the Thesis

The objectives of this thesis are divided into two parts, corresponding to the theoretical and practical aspects of the work.

The theoretical part aims to provide a comprehensive survey of the challenges associated with project adaptation in the context of full line code completion, focusing on the current capabilities and limitations of in-context learning and repository-level code completion methods. This part establishes the necessary background to understand the practical contributions of the thesis.

The practical part addresses several research questions concerning the impact of various context composition strategies on the quality of code completion, evaluated across three setups: a pre-trained Code Large Language Model (Code LLM), the same model fine-tuned with a specific context composition strategy, and the model after context window extension. To address these questions, the thesis implements a context composition framework for extracting relevant information from software repositories, as well as a fine-tuning pipeline for adapting Code LLMs to project-specific data. The outcomes characterize the role of context composition in enhancing code completion quality and provide insights into its integration within repository-level training procedures.

# Introduction

Software development plays a pivotal role in shaping the modern digital world. Assisting developers in writing code more efficiently has the potential to accelerate innovation across industries by reducing the cognitive and temporal overhead of programming. Over the decades, numerous tools have emerged to support developers, including high-level programming languages, integrated development environments (IDEs), version control systems, and, more recently, AI-powered code assistance. Tasks such as code completion, generation, refactoring, and bug localization are increasingly addressed using Large Language Models (LLMs), which have demonstrated strong capabilities in natural language and code understanding.

Code completion, in particular, benefits significantly from the in-context learning capabilities of LLMs — where the model leverages examples or relevant information provided in the input context to improve predictions without parameter updates. However, while LLMs have advanced in modeling local contexts, a critical limitation remains: their restricted ability to integrate and reason over information dispersed across large codebases. This includes understanding dependencies between files, class hierarchies, and interactions with external libraries — information that is often essential for generating coherent and accurate completions.

Despite the emergence of code-specific LLMs and efforts to incorporate repository-level context during training or inference, most pre-trained models still struggle to process more than a small fraction of a project’s code files at once. As a result, they fall short in capturing project-wide structure, leading to incomplete or inaccurate predictions. Therefore, there is a need for continued research within the community developing such models to address these limitations. This thesis focuses on exploring how the composition and organization of contextual information influence the performance of code completion models, particularly in repository-level settings where code is distributed across multiple interdependent files. It investigates how context selection strategies and model adaptations can enhance the ability of pre-trained Code LLMs to generate coherent and accurate completions that reflect the structure and se-

mantics of entire software projects.

The thesis is organized as follows: Chapter 1 offers an overview of the code completion task, including its definition, significance, historical approaches, and taxonomy. Chapter 2 introduces essential concepts related to language modeling (LM) and state-of-the-art architecture for this domain within machine learning. In Chapter 3, the discussion on language modeling is extended to encompass code completion with relevant details. Chapter 4 covers the aspect of in-context learning, emphasizing its relevance to repository-level capabilities. Chapter 5 provides an overview of the primary issues associated with long contexts and the methods devised to address them. The implementation of the tools used to answer the research questions is elaborated in Chapter 6. Finally, Chapter 7 describes the research, including precise question formulations, experimental design, and results interpretation.

Part I

Conceptual Framework

# Chapter 1

## Code Completion

The upcoming chapter provides a comprehensive examination of code completion. It begins with a definition of the task. The chapter then delves into the motivations for advancing code completion techniques, highlighting their impact on productivity and learning. A historical overview of the evolution of code completion methods is presented, tracing the transition from deterministic approaches to sophisticated learning-based systems. Finally, the chapter categorizes code completion based on various criteria.

### 1.1 Task Definition

Code completion, also called code suggestion or autocompletion, is a functionality within integrated development environments that enables developers to utilize an automated continuation of the code they are writing. Typically, it is invoked by a trigger model, which aims to anticipate the user's requirement for this functionality. The point of this invocation, whether it is referenced as temporal or spatial, is called a trigger point.

In this work, code completion is regarded as a task rather than a feature, as this perspective highlights the task-methodology relationship and emphasizes the technical aspects over marketing considerations. In addition, the term *completion* is used as a shorthand throughout the text.

### 1.2 Motivation

The task of code completion is a pivotal component of contemporary software development, offering substantial benefits that enhance both the productivity and efficiency of developers. This section elucidates the motivations for investigating this task and advancing existing solutions.

By reducing the amount of typing required, code completion allows developers to focus more on the logic and structure of their code rather than the



syntax. Studies have shown that AI-powered code completion tools can significantly reduce task execution times, with some reporting up to a 55.8% reduction in controlled experiments (Peng et al. 2023). These tools also significantly impact developer productivity by reducing the cognitive load associated with coding tasks, achieved through contextually relevant suggestions that streamline the coding process (Weber et al. 2024). Furthermore, the integration of AI in code completion tools has been linked to increased developer satisfaction and efficiency, as it allows developers to complete tasks more swiftly and with greater accuracy (Bakal et al. 2025).

For students and novice programmers, code completion tools serve as a valuable learning aid. They provide real-time feedback and suggestions, helping learners understand programming concepts and syntax more effectively. This educational aspect is highlighted in studies where students reported enhanced learning experiences and increased confidence in their coding abilities (Takerngsaksiri et al. 2023). Moreover, code completion tools provide accurate suggestions that help reduce typo errors and other common mistakes. This is particularly beneficial for new developers or those working with unfamiliar codebases, as it helps them adhere to coding standards and best practices.

Recent advancements in deep learning have enabled the personalization of code completion tools, allowing them to adapt to specific organizational or individual coding styles. This personalization not only improves the relevance of suggestions but also enhances the overall user experience, making these tools more effective and user-friendly (Giagnorio et al. 2025).

Due to the simple and atomic formulation of code completion, it possesses the important property of being fundamental to other AI-powered code assistance features. For instance, code generation can be viewed as a specific extension of code completion, and code editing is a sequence of code generations. This implies that most enhancements achieved through code completion research propagate further, adding new layers of improvements in the field.

In summary, code completion stands as a foundational and multifaceted tool within modern software development. Its benefits extend beyond mere convenience, offering measurable improvements in efficiency, learning, and code quality. As advancements in AI continue to refine its capabilities, code completion not only streamlines day-to-day programming tasks but also acts as a cornerstone for broader innovations in AI-assisted software engineering.

### 1.3 Evolution of Methods

Approaches to code completion have evolved considerably over the past few decades, progressing from simple deterministic methods to sophisticated learning-based systems. This section highlights key developments in the history of code completion. The Related Work section of the Ciniselli et al. (2021) was mainly used to compile this valuable information.

In the early years, code completion methods primarily relied on rule-based

approaches and static type information. These systems typically presented the user with a list of type-compatible methods or variables, often sorted alphabetically (Mandelin et al. 2005). While straightforward to implement, these methods lacked contextual awareness and often produced lengthy suggestion lists that were cumbersome to navigate.

A significant advancement came with clone-based completion methods, where code fragments from existing repositories were identified and reused as completion suggestions (Hill et al. 2004). These approaches recognized that developers often write repetitive code patterns, but they were limited by the need for exact or near-exact matches in the code database.

The next major evolution occurred when researchers began applying statistical methods to code completion. Han et al. (2009) introduced a novel approach that expanded abbreviated inputs into complete code tokens using Hidden Markov Models trained on example code. This represented an early step toward probabilistic code completion. Hindle et al. (2012) demonstrated that software exhibits natural patterns that can be captured by statistical language models, laying the groundwork for treating code completion as a language modeling problem.

Building on this foundation, more sophisticated probabilistic models were developed. Raychev et al. (2014) applied statistical language models specifically for code completion tasks, while Bielik et al. (2016) introduced probabilistic higher order grammar (PHOG), which parameterized grammar production rules on a context obtained from executing functions learned from data. Bayesian networks were also explored by Proksch et al. (2015), who used additional context information beyond just the static type to generate completion suggestions.

Context-sensitivity became increasingly important as code completion systems matured. Asaduzzaman et al. (2014) developed CSCC, which leveraged previous code examples to recommend method calls by considering the surrounding code context. This approach demonstrated how incorporating more contextual information affects the relevance of completion suggestions.

The rise of deep learning brought another paradigm shift to code completion. Recurrent neural networks (RNNs), particularly long short-term memory (LSTM) networks, were applied to model the sequential nature of code. Ginzberg et al. (2017) explored both standard LSTMs and attention-augmented networks for token-level code completion, establishing new capabilities in code prediction tasks.

Recent years have witnessed the adoption of even more powerful neural architectures. Karampatsis et al. (2019) tackled the out-of-vocabulary problem in code completion using subword units with neural language models. Liu et al. (2020) leveraged multi-task learning and pre-trained language models for code completion, while also incorporating type information to assist in identifier prediction.

The latest advancements utilize transformer-based architectures, which

have proven effective in natural language processing (NLP). Svyatkovskiy et al. (2020) introduced IntelliCode Compose, a general-purpose code completion tool that predicts sequences of code tokens using a transformer model. Kim et al. (2021) enhanced transformer models by incorporating syntactic structure awareness through abstract syntax tree (AST)-based representations. These transformer architectures address limitations of previous sequence models by allowing attention to any part of the input context, rather than relying solely on sequential information. Furthermore, attention mechanisms enable transformers to consider relationships between distant parts of the code that have semantic connections, overcoming the long-range dependency challenges faced by RNN-based approaches.

Parallel to these developments, structural approaches that go beyond treating code as a mere sequence of tokens have gained prominence. Alon et al. (2019) proposed a structural language modeling approach that leverages the strict syntax of programming languages to model code as a tree, capturing both syntactic and semantic relationships in the code. These structure-aware models offer an alternative to purely sequential approaches by explicitly modeling the hierarchical nature of source code.

The evolution of code completion methods reflects a general trend toward more contextually aware, semantically rich, and structurally informed models that capture the unique characteristics of source code compared to natural language. It also emphasizes the dominance of the probabilistic approaches over the deterministic ones.

## 1.4 Taxonomy of Code Completion

The concept of code completion lacks a strict definition within the field, as its interpretation has evolved with the various approaches employed to address it over time (see Section 1.3 for more details). This section provides a comprehensive overview of the diverse characteristics of code completion and highlights the specific focus of this thesis.

To grasp the meaning of the subsequent text, it is crucial to understand the term *token*, which refers to any subsequence of characters with a finite length.

### Granularity

Completion can be applied under the different degrees of granularity, which refers to the scope and detail of the produced code. At the most basic level, next-token completion involves predicting the subsequent token, such as a keyword, operator, or identifier. This level of granularity is useful for fine-grained suggestions that assist developers in writing code efficiently. Moving up in complexity, single line completion involves predicting the continuation of the given line of code based on the current context, which requires a broader understanding of the code's structure and logic. At the highest level, code

block completion involves generating entire blocks of code, such as functions or classes, which necessitates a comprehensive understanding of the codebase and its architecture. This level of granularity is akin to code generation, where the model not only completes existing code but also creates new, coherent code structures. Each level of granularity serves different purposes and can be leveraged depending on the specific needs of the development task at hand.

For this work, single-line completion is selected because it presents a greater challenge for modern models compared to next-token prediction and serves as a fundamental baseline task for evaluating the performance of the proposed methods.

## Context

In the realm of code completion, a distinction is drawn between file-level and repository-level approaches, each characterized by its scope and contextual depth. File-level code completion operates within the confines of a single file, leveraging the immediate context such as local variable declarations, function definitions, and imports. This approach is effective for small, isolated scripts but is limited in its ability to capture the broader interactions present in larger projects. Conversely, repository-level code completion extends its reach to encompass the entire codebase, integrating information from multiple files and modules. This broader context is essential for understanding complex dependencies and interactions that span across the repository, enabling more accurate and contextually relevant code suggestions. The need for greater context in repository-level completion is driven by the intricate and interconnected nature of modern software systems, where a comprehensive understanding of the entire codebase is crucial for effective code completion.

Repository-level completion has been selected as the focus of this thesis because it offers a broader research landscape and addresses a significant need within the field.

## Suffix Awareness

The completion scenario can be categorized based on the directionality of the prediction relative to the trigger point. The left-to-right scenario involves generating code predictions using only the context preceding the trigger point, which is typical in traditional code completion systems. This approach is effective for straightforward code continuations but may lack the ability to consider the broader context. Conversely, the left-and-right scenario, also known as bidirectional completion, leverages both the preceding and succeeding context around the trigger point. This method allows for more informed predictions by considering the entire file, thus enhancing the accuracy and relevance of the suggestions.

Modern models for code completion employ the fill-in-the-middle (FIM) approach to gain suffix awareness capability. Although this thesis describes

this method in Section 3.3, it does not expand on it and instead focuses on the more fundamental left-to-right scenario. This choice is driven by the practical part of this work and its design constraints, which are motivated in Section X.

**TODO:** “Add a forward reference to the experiment design.”

### Line Start Availability

Another criterion for categorizing completion systems is whether the initial portion of the target line has already been completed by the user. In practical applications, this scenario is quite common. However, for the purpose of model evaluation, it is more straightforward to assess the full line completion, where the initial part of the target line is not provided. This approach serves as a lower bound for assessing completion quality, as the ability to predict the entire line demonstrates a stronger capability.

Due to the aforementioned reasoning, the full line code completion is further considered throughout this thesis.

### Programming Language Support

The code completion systems can support either a single programming language (PL) or multiple languages simultaneously.

In the practical part of this thesis, multilingual models are employed with a focus on a single PL, specifically Python<sup>1</sup>. This design choice is motivated by resource constraints.

To finalize the terminology used in subsequent parts of this thesis regarding this task, the *completion* refers to a full, single-line, repository-level code completion without suffix awareness.

---

<sup>1</sup><https://www.python.org/>

## Chapter 2

# Standard LM

This chapter provides the necessary theoretical foundation to understand general language modeling, without specifically focusing on the code completion task. It begins with the formulation and presents a probabilistic perspective on the models employed to represent language. Subsequently, it explores the transformer architecture and the training of such models. The chapter concludes with a discussion on the inference and sampling processes of trained models.

In this chapter, the term *language* refers to an arbitrary natural language, as opposed to a formal one.

### 2.1 Definition

Consider a natural language denoted as  $\mathcal{L}$ , which is a set of all possible sentences in this language. A sentence  $\mathbf{x}_{1:T}$  of length  $T$  is a sequence of words  $(x_1, x_2, \dots, x_T)$ , where each word is an element of the vocabulary  $V$ . Given that some sentences are more likely to occur than others, there exists a joint probability distribution  $P : \mathcal{L} \rightarrow [0, 1]$  over variable-length sequences from  $\mathcal{L}$ . The model  $p_\theta$  with parameters  $\theta$  (also called weights) that represents this distribution is called a language model, and the task of creating such a model is referred to as language modeling.

Although natural languages are infinite due to their recursive structure, they have a finite number of words, meaning  $K = |V| \in \mathbb{N}$ . To demonstrate this, consider that word lengths are finite, which implies that  $\mathcal{L}$  certainly contains a word with the maximum length  $M$ . Let the number of characters involved in word formation be  $N$ . A very rough upper bound estimate for the number of words in  $\mathcal{L}$  can be calculated using the formula  $N^M$ , which is finite because both  $N$  and  $M$  are finite.

The chain rule of probability can be applied to represent  $P$  as follows:

$$P(\mathbf{x}_{1:T}) = P(x_1) \cdot P(x_2 | x_1) \cdot P(x_3 | x_2, x_1) \cdots = \prod_{t=1}^T P(x_t | \mathbf{x}_{1:t-1}) \quad (2.1)$$

This indicates that language modeling is equivalent to estimating the probability of each word in the sentence given all the preceding words (i.e., the context), where the probability distribution over each word is a  $K$ -dimensional categorical distribution.

This indicates that language modeling is equivalent to estimating the probability of each word in the sentence given all the preceding words (i.e., the context), where the probability distribution over each word is a  $K$ -dimensional categorical distribution. This sequential prediction process, where each word depends only on the prior subsequence, is called autoregressive (AR) language modeling.

## 2.2 Text Representation

The proper representation of text is critical to effective language modeling design. Treating text as a sequence of words is natural and intuitive for humans, but not the best choice for machines.

### 2.2.1 Tokens

The first reason why words may not be the optimal choice for representing discrete units of text is the out-of-vocabulary (OOV) problem, which refers to the presence of words (or tokens) that are absent from a language model's vocabulary. This challenge arises from the inherently dynamic nature of language, which lacks fixed boundaries. As language evolves, some words become obsolete while new ones emerge. Additionally, the text corpora used for training language models often contain typographical errors and other artifacts of human language.

The second reason is that the language models require statistical information about the co-occurrence of words to capture the semantics of the language. It becomes a problem with rare words whose representations in models lack a learning signal. This phenomenon can be lucidly demonstrated with Zipf's law, which states that in a given corpus, the frequency of any word is inversely proportional to its rank in the frequency table (Estoup 1916). This implies that a small number of words are used very frequently, while the majority are rare. Consequently, language models trained on word-level representations may struggle to learn these infrequent words, leading to poor generalization.

One can propose employing character-level text representation. Although this approach completely resolves the OOV problem, it introduces its own limitations. A small vocabulary necessitates processing longer sequences with

smaller semantic representation of its units, which introduces additional computational overhead to internally combine these units to form a more integral semantic representation of text.

Thus, there is a trade-off between the size of the vocabulary and the size of input sequences. Minimizing one leads to an increase in the other. A large vocabulary means that some of its units are used very rarely, and their representations lack a learning signal. A small vocabulary means the utilization of more semantically poor units and a greater computational overhead to process longer sequences.

The optimal balance between character-level and word-level representations is subword chunks, generally called tokens. They offer a greater prospect of granularity in text division, which allows addressing the aforementioned trade-off on a more fine-grained level.

### 2.2.2 Byte Pair Encoding for Tokenization

The OOV problem persists at the token level. A variety of methods exist to address this issue, with one of the most popular being the byte pair encoding (BPE) algorithm. Originally proposed by Gage (1994) as a compression algorithm, it was later adapted by Sennrich et al. (2015) for word segmentation with the aim of constructing a vocabulary.

The BPE process begins by initializing the vocabulary with individual characters and a special end-of-word symbol. Words are initially represented as sequences of these characters. Given some corpus of text, the algorithm iteratively identifies the most frequent pair of adjacent symbols and merges them into a new symbol. This process continues until a predefined number of merge operations is reached, resulting in a vocabulary composed of both characters and frequently occurring subword units. After the vocabulary is constructed, the same algorithm can be applied to tokenize the text.

This approach allows for a compact representation of text, reducing the sequence length while maintaining the ability to represent rare and unseen words. By balancing the trade-off between vocabulary size and sequence length, BPE enables language models to efficiently handle the dynamic nature of language, capturing both common and rare linguistic patterns.

### 2.2.3 Embeddings

After the text is tokenized, it is represented as a sequence of tokens  $(x_1, x_2, \dots, x_T)$ , where  $x_i \in \{1, 2, \dots, K\}$  is a corresponding integer identifier. A parameterized language model  $p_\theta$  should have the ability to process this sequence of integers. Since most effective models are based on neural networks (NNs),  $\theta$  is a set of matrices that define some non-linear transformation over the continuous vector space. Therefore, tokens need to be embedded into the input vector space  $\mathbb{R}^d$ .



The most straightforward approach is to set  $d = K$  and use one-hot encoding, which assigns a vector  $\mathbf{x}_i = \text{one-hot}(x_i) \in \{0, 1\}^K$  full of zeros except for the position corresponding to the token, which is set to one. The resulting vectors capture the identity of the tokens but lack the interaction effects between them.

This issue is resolved by incorporating the first layer of the NN as a trainable embedding matrix  $\mathbf{E} \in \mathbb{R}^{d' \times K}$ , which maps each token to a vector  $\mathbf{E}\mathbf{x}_i$  in  $\mathbb{R}^{d'}$ , called an embedding, by multiplying its one-hot encoding with  $\mathbf{E}$ . Note that the sparse nature of one-hot representations allows this multiplication to be implemented using a lookup table by taking the  $i$ -th column of  $\mathbf{E}$ .

There are multiple ways to train the embedding matrix  $\mathbf{E}$ . The choice of method depends on the desired properties. In the case of language modeling, the primary requirement is the alignment of the embeddings with the rest of the model weights. Therefore, the embeddings are frequently initialized and trained jointly with the rest of the model parameters.

The trained embeddings capture the semantics of the tokens. This means that tokens occurring in similar contexts tend to have similar embeddings. This phenomenon is known as the distributional hypothesis (Harris 1954). The most common measure of similarity between two embeddings is their dot product or its normalized version, cosine similarity:

$$\cos(\mathbf{x}, \mathbf{y}) = \frac{\mathbf{x}^\top \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|}. \quad (2.2)$$

As the dot product suffers from the curse of dimensionality, the number of dimensions  $d'$  required to capture all semantic patterns is exponentially smaller than the vocabulary size  $K$ . This is achieved by the fact that two dissimilar vectors do not necessarily have to be orthogonal, as their dot product is already near zero.

## 2.3 Transformer Models

The most widely used architecture for language modeling is the Transformer model proposed by Vaswani et al. (2017). It is based on the attention mechanism, with its roots introduced by Bahdanau et al. (2014). Since then, a vast number of various architectural modifications have appeared. However, the core idea remains the same.

In the primary source paper, the authors proposed two parts of the architecture: the encoder and the decoder. Both of them, or their combination, are used depending on the specific task. In the case of language modeling (including code modeling), the decoder-only variant has gained widespread adoption. This part of the text describes only this type of Transformer.

### 2.3.1 Overview of the Forward Pass

Consider the sequence of context tokens  $\mathbf{x}_{1:t-1} = (x_1, x_2, \dots, x_{t-1})$  from equation 2.1. To model the probability  $p_\theta(x_t \mid \mathbf{x}_{1:t-1})$  of the next token  $x_t$ , the decoder-only Transformer first embeds the context tokens into vectors  $(\mathbf{Ex}_1, \mathbf{Ex}_2, \dots, \mathbf{Ex}_{t-1})$  and then passes them through the stack of decoder layers to obtain the contextually enriched embeddings  $(\mathbf{h}_1, \mathbf{h}_2, \dots, \mathbf{h}_{t-1})$ , which are called hidden states. It then takes the last hidden state and multiplies it by the weights matrix to get the logits  $\mathbf{z}_t = \mathbf{W}\mathbf{h}_{t-1}$ . The  $\mathbf{W}$  is often called the language modeling head. The softmax function is then used to get the parameters of the categorical distribution over the vocabulary:

$$\mathbf{p}_t = \text{softmax}(\mathbf{z}_t) = \frac{\exp(\mathbf{z}_t/\tau)}{\sum_{k=1}^K \exp(z_{t,k}/\tau)}, \quad (2.3)$$

where the exponent in the numerator and division are element-wise functions, and  $\tau$  is a temperature parameter discussed in Section 2.5.1.

The resulting vector of estimated probabilities  $\mathbf{p}_t$  can be used in multiple ways. First, to fulfill the original task of language modeling, the element corresponding to the next token  $x_t$  can be extracted from  $\mathbf{p}_t$  and serve as  $p_\theta(x_t \mid \mathbf{x}_{1:t-1})$  in equation 2.1. Second, one can sample from the categorical distribution defined by  $\mathbf{p}_t$  to get the next token and thus generate a new sequence from the model, which means that the language model can be used as a generative model. Both settings are commonly used. The first one is employed for the training process, while the second one is used for inference.

Before delving into the specifics of the decoder layer, two important points should be noted. First is that the initial token lacks context. Consequently, a special beginning of sequence (BOS) token is utilized to provide context and serve as an embedding for  $p_\theta(x_1)$ . Second, for simplicity, the bias term is omitted in all the weight application formulas, as its inclusion is optional and contingent upon the engineer's design choices.

### 2.3.2 Decoder Layer

Let's denote the input embeddings as  $(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T)$ , where  $\mathbf{x}_i \in \mathbb{R}^{d_{\text{model}}}$  for each  $i \in \{1, 2, \dots, T\}$ , and let  $\mathbf{X} \in \mathbb{R}^{T \times d_{\text{model}}}$  represent their concatenated form. The decoder layer then applies the following operations:

$$\mathbf{X}^{(1)} = \mathbf{X} + \text{MultiHead}(\text{Norm}(\mathbf{X})), \quad (2.4)$$

$$\mathbf{X}^{(2)} = \mathbf{X}^{(1)} + \text{MLP}(\text{Norm}(\mathbf{X}^{(1)})), \quad (2.5)$$

where  $\mathbf{X}^{(2)} \in \mathbb{R}^{T \times d_{\text{model}}}$  is the output of the layer, and Norm, MultiHead, MLP are sub-layers discussed below.

These formulas differ from the original architecture described in Vaswani et al. (2017) as they employ pre-normalization (Baevski et al. 2018; Child et al.

2019; Wang et al. 2019). This modification enhances the stability of the training process by ensuring that the residual embedding stream is modified only through addition, which helps maintain well-behaved gradients at initialization (Xiong et al. 2020).

### 2.3.3 Normalization

The normalization layer stabilizes the training process by normalizing the input embeddings independently for each token. Various types of normalization layers exist. While the original architecture used layer normalization (Ba et al. 2016), root mean square layer normalization (Zhang et al. 2019) has also been widely adopted.

### 2.3.4 Attention

Consider two additional architectural parameters  $d_k$  and  $d_v = d_{\text{model}}$ . The attention layer is described using three weight matrices:  $\mathbf{W}_Q, \mathbf{W}_K \in \mathbb{R}^{d_k \times d_{\text{model}}}$  and  $\mathbf{W}_V \in \mathbb{R}^{d_v \times d_{\text{model}}}$ . The output of the layer is calculated as follows:

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax} \left( \frac{[\mathbf{Q}\mathbf{K}^\top]_{\text{mask}}}{\sqrt{d_k}} \right) \mathbf{V}, \quad (2.6)$$

where  $\mathbf{V} = \mathbf{X}\mathbf{W}_V^\top$  and  $\{\mathbf{Q}, \mathbf{K}\} = \text{RoPE}(\mathbf{X}\mathbf{W}_{\{Q,K\}}^\top)$  with RoPE denoting the rotary position embedding discussed in Section 2.3.8. The softmax is applied row-wise, and  $[\cdot]_{\text{mask}}$  is a causal masking operation that sets all elements above the main diagonal to  $-\infty$ , thereby preventing data leakage from future tokens.

Attention functions as a layer that enables tokens to communicate with each other. Each token performs a query  $\mathbf{q}_i$  to all preceding tokens and itself  $\mathbf{k}_{j \leq i}$ . The corresponding attention weight of the query-key match  $\mathbf{q}_i^\top \mathbf{k}_j$  is normalized by the softmax and used as a soft dictionary lookup to retrieve the partial component of the value vector  $\mathbf{v}_j$ . The convex combination of all such value vectors  $\mathbf{v}_{j \leq i}$  becomes the output of the attention layer.

### 2.3.5 Multi-Head Attention

Multi-head attention is a modification of the basic attention mechanism. Instead of applying the attention mechanism once, the embedding processing is divided into multiple heads that perform attention in parallel. Given the number of heads  $h$  and an additional output projection weight matrix  $\mathbf{W}_O \in$

$\mathbb{R}^{hd_v \times d_{\text{model}}}$ , the output of the layer is calculated as:

$$\text{MultiHead}(\mathbf{X}) = \quad (2.7)$$

$$\text{MultiHead}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{Concat}(\text{head}_1, \dots, \text{head}_h) \mathbf{W}_O \quad (2.8)$$

$$\text{where head}_i = \text{Attention}(\mathbf{Q}_i, \mathbf{K}_i, \mathbf{V}_i), \quad (2.9)$$

where  $\mathbf{Q}_i, \mathbf{K}_i, \mathbf{V}_i$  are obtained by applying an independent set of weight matrices  $\mathbf{W}_Q^{(i)}, \mathbf{W}_K^{(i)}, \mathbf{W}_V^{(i)}$  to the input embeddings  $\mathbf{X}$ , and  $\mathbf{Q}, \mathbf{K}, \mathbf{V}$  are their concatenated versions.

In this version of the attention,  $d_v$  is not necessarily equal to  $d_{\text{model}}$ .

### 2.3.6 Multi-Layer Perceptron

A multi-layer perceptron (MLP) is a simple feed-forward neural network with two linear transformations and a non-linear activation function. It is applied to each token independently.

Originally, the rectified linear unit (ReLU) was used as the activation function. However, other alternatives have empirically been shown to perform better. Two of these are the Gaussian error linear unit (GELU) and the sigmoid linear unit (SiLU) from Hendrycks et al. (2016). Another is the swish gated linear unit (SwiGLU) introduced in Shazeer (2020).

### 2.3.7 Dropout

Dropout is a widely adopted regularization technique that enhances the generalization of the model (Srivastava et al. 2014). During the training process, it randomly masks some of the activations of the preceding layer with a probability  $p_{\text{dropout}}$  and scales the remaining activations by a factor of  $1/p_{\text{dropout}}$  to maintain the output statistics of the layer.

This method is frequently applied in transformers across all layers, including the dropout of the attention scores calculated before the application of the softmax function.

### 2.3.8 Rotary Position Embedding

Rotary Position Embedding (RoPE) is a significant architectural enhancement that warrants its own section. It is a modification built upon the vanilla Transformer to incorporate the positional information of tokens and address issues present in the originally proposed analogous approach. RoPE was introduced by Su et al. (2021) and quickly gained widespread adoption.

For each embedding  $\mathbf{x}_m$  at position  $m$  in the sequence, RoPE proposes a rotation of its query and key vectors:

$$\{\mathbf{q}_m, \mathbf{k}_m\} = \text{RoPE}(\mathbf{W}_{\{Q,K\}} \mathbf{x}_m, m) = \mathbf{R}_{\Theta, m}^d \mathbf{W}_{\{Q,K\}} \mathbf{x}_m \quad (2.10)$$

using an orthogonal rotary matrix

$$\mathbf{R}_{\Theta, m}^d = \begin{pmatrix} \cos m\theta_1 & -\sin m\theta_1 & 0 & 0 & \cdots & 0 & 0 \\ \sin m\theta_1 & \cos m\theta_1 & 0 & 0 & \cdots & 0 & 0 \\ 0 & 0 & \cos m\theta_2 & -\sin m\theta_2 & \cdots & 0 & 0 \\ 0 & 0 & \sin m\theta_2 & \cos m\theta_2 & \cdots & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & \cdots & \cos m\theta_{d/2} & -\sin m\theta_{d/2} \\ 0 & 0 & 0 & 0 & \cdots & \sin m\theta_{d/2} & \cos m\theta_{d/2} \end{pmatrix}, \quad (2.11)$$

where  $d = d_k$  and  $\Theta = \{\theta_i = \theta_{\text{base}}^{-2(i-1)/d}, i \in \{1, 2, \dots, d/2\}\}$  is a predefined set of frequencies derived from the base frequency hyperparameter  $\theta_{\text{base}}$ .

Note two important details regarding the application of RoPE. First, it requires  $d$  to be even. Second, it is applied across all heads in multi-head attention without repeating the rotation frequency  $h$  times, inherently adding more functional bias to heads working on higher frequencies (Barbero et al. 2024).

When multiplying query and key matrices in equation 2.6, the two embeddings  $\mathbf{x}_m$  and  $\mathbf{x}_n$  derive the attention score decay based on the relative distance  $m - n$  between them:

$$\mathbf{q}_m^\top \mathbf{k}_n = (\mathbf{R}_{\Theta, m}^d \mathbf{W}_Q \mathbf{x}_m)^\top (\mathbf{R}_{\Theta, n}^d \mathbf{W}_K \mathbf{x}_n) = \mathbf{x}_m^\top \mathbf{W}_Q \mathbf{R}_{\Theta, n-m}^d \mathbf{W}_K \mathbf{x}_n, \quad (2.12)$$

which allows the model to capture the positional information of the tokens.

## 2.4 Training

### 2.4.1 Maximum Likelihood Estimation

Training the language model  $p_\theta$  involves aligning the distribution it represents with the true data distribution. However, the later is unknown, and only a finite set  $\mathcal{D} = \{\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(n)}\}$  of instances sampled from it is available. The training objective  $\mathcal{L}_{\mathcal{D}}$ , known as the likelihood function, is expressed as the probability of these samples being generated by the model. Assuming that samples are independent and identically distributed,  $\mathcal{L}_{\mathcal{D}}$  can be decomposed into the product of individual probabilities:

$$\mathcal{L}_{\mathcal{D}}(\theta) = p_\theta(\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(n)}) = \prod_{i=1}^n p_\theta(\mathbf{x}^{(i)}) \rightarrow \max. \quad (2.13)$$

The method of estimating the parameters  $\theta$  by maximizing this function is called maximum likelihood estimation (MLE).

Due to several reasons, including underflow issues, the log-likelihood function is often used instead:

$$\ell_{\mathcal{D}}(\theta) = \log \mathcal{L}_{\mathcal{D}}(\theta) = \log \prod_{i=1}^n p_\theta(\mathbf{x}^{(i)}) = \sum_{i=1}^n \log p_\theta(\mathbf{x}^{(i)}). \quad (2.14)$$

Optimizing both functions is equivalent, as the logarithm is a monotonically increasing function.

To take one step further, a minus sign can be added to focus on minimization instead of maximization, thereby formulating the negative log-likelihood loss function. Additionally, equation 2.1 can be used to express the loss function as a sum of the log-likelihoods of the individual tokens:

$$L(\theta) = -\ell_{\mathcal{D}}(\theta) = -\sum_{i=1}^n \sum_{t=1}^{T_i} \log p_{\theta}(x_t^{(i)} \mid \mathbf{x}_{1:t-1}^{(i)}) \rightarrow \min. \quad (2.15)$$

This objective corresponds to the cross-entropy (CE) loss, which is often represented using the dot product of  $\mathbf{p}_t^{(i)}$  from equation 2.3 and the one-hot encoding of the target token  $x_t^{(i)}$ .

### 2.4.2 Gradient-Based Optimization

Various optimization algorithms have been developed to train neural networks. The most effective ones are based on first-order methods, which iteratively adjust the parameters by subtracting the stochastic gradient of the loss function with respect to the parameters. The degree of this adjustment is controlled by multiplicative factor of the gradients called learning rate. Gradients are calculated using a subsample of the training set, known as a mini-batch (or simply batch), with the application of automatic differentiation.

In the field of language modeling, the adaptive moment estimation (Adam) (Kingma et al. 2014), a modification of stochastic gradient descent (SGD), and its version supporting weight decay (Loshchilov et al. 2017) are the most popular algorithms. Adam maintains two moving average statistics, called the first and second moments, for each parameter, which correct the application of the gradient and accelerate convergence.

### 2.4.3 Batch-Related Terminology

Due to the operational nature of graphics processing units (GPUs), multiple sequences must be organized to have the same length to be processed in parallel. This is achieved by either padding shorter sequences with special tokens or truncating longer ones. The model's scope of visibility is termed the context window, and its size is referred to as the context length.

In some cases, sequences are randomly shuffled, concatenated, and split to form batches without the need for padding or truncation. This technique is known as batch packing.

When hardware constraints limit batch size, gradient accumulation is utilized. Instead of using the entire batch for a single optimization step, it is divided into micro-batches, which are processed sequentially. The parameters are updated only after all the micro-batches have been processed.

### 2.4.4 Learning Rate Scheduling

Due to the design of Adam, a few iterations at the beginning of training are required to gather sufficient gradient statistics for better moment estimation. This requirement becomes even more critical in the context of the instability associated with Transformers. During this initial phase, using the full learning rate is inadvisable, as it may lead to unstable updates. To mitigate this issue while utilizing all available data, a warm-up period is introduced in the training process. Throughout this period, the learning rate is gradually increased, often linearly, from zero to the desired value over a predefined number of iterations.

After training with high learning rates, the model benefits from smaller parameter updates as it approaches the local optimum and requires more precise adjustments. This is achieved by decaying the learning rate over time.

To summarize, the learning rate, rather than being a constant hyperparameter, has evolved into a function of the iteration index, thereby opening a new dimension for controlling the training process.

## 2.5 Inference

### 2.5.1 Sampling

As previously mentioned, a language model represents a probability distribution over the vocabulary given the preceding context. The iterative process of sampling from this distribution and incorporating these new tokens into the model's context is referred to as generation.

Tokens are selected based on their associated probabilities  $\mathbf{p}_t$  from the softmax output, as described in equation 2.3. The temperature parameter modulates the sharpness of this distribution. There are two limiting cases: First, as  $\tau \rightarrow +\infty$ , the distribution becomes uniform, with all tokens having an equal probability of being selected. Second, as  $\tau \rightarrow 0^+$ , the selection becomes deterministic, choosing the most probable token. This is known as greedy decoding and corresponds to the greedy search algorithm in a graph of all possible sequences, where the language model functions as a heuristic.

To further enhance the quality of the generated text, the top- $k$  sampling strategy is often employed. It restricts the vocabulary to the top- $k$  most probable tokens, which helps prevent the model from outputting irrelevant tokens.

A similar concept is applied in the nucleus sampling strategy, also known as top- $p$  sampling (Holtzman et al. 2019). It restricts the vocabulary to the smallest set of tokens whose cumulative probability exceeds a threshold parameter  $p$ .

### 2.5.2 Stopping Criteria

There are several methods to terminate the generation process.

Firstly, the model can signal the completion of the output sequence by generating an end of sequence (EOS) token. This requires the EOS token to be included in the vocabulary, and the model must be trained to produce it. This is accomplished by appending the EOS token to the end training sequences.

Secondly, algorithmic stopping criteria can be employed. For example, the generation process can be halted when a maximum length threshold is exceeded. Another pertinent example, particularly in the context of code completion, is to cease generation when the model outputs a token containing a newline character.



# Completion-Centric LM

TODO: “Description of the chapter.”

## 3.1 Training Stages

The training of modern Code LLMs begins with acquiring a model checkpoint that possesses a robust understanding of code. This can be achieved through two distinct types of training stages. One type involves taking a general-purpose pre-trained LLM and fine-tuning it on a file-level code dataset. This approach is utilized by models such as Code Llama (Rozière et al. 2023) and Qwen2.5-Coder (Hui et al. 2024). The other type involves training a model from scratch using a mixture of code and code-related natural language data. This pre-training approach is adopted by models like DeepSeek-Coder (Guo et al. 2024) and OpenCoder (Huang et al. 2024). Both types necessitate vast amounts of data, typically ranging from 2 trillion (2T) to 18 trillion (18T) tokens with context length of 4,096 (4K).

Once a strong model with a limited context size is trained, it often undergoes repository-level pre-training (sometimes referred to as long-context fine-tuning). The objective of this stage is to extend the context window, thereby enabling the model to comprehend a broader scope within a given repository. Unlike the initial training stage, repository-level pre-training requires significantly fewer tokens. However, it necessitates the application of context extrapolation methods and the utilization of longer input sequences. This phase typically involves between 8 billion (8B) and 300 billion (300B) training tokens, with the context window extended with 16K or 32K tokens.

Each input sequence in repository-level pre-training consists of two components: the composed context and the completion file. Various methods exist for obtaining the former. In this thesis, the function responsible for this task is referred to as the context composer, or simply composer. This process involves retrieving and preprocessing a subset of files, which are then assembled

to form a repository context.

Models trained using only the aforementioned stages are referred to as base models. To enhance their utility for various tasks, an additional instruction tuning phase is often conducted. However, the capabilities gained from this stage are not essential for the code completion task and are not discussed further in this thesis. Throughout this work, all mentioned models are assumed to be their base versions.

## 3.2 Gradient Masking

The learning signal for the model is derived from both the composed context and the completion file sources. The alignment of their distributions is contingent upon the specific choice of the composer. When learning the completion of all tokens, a mismatch in these distributions can introduce undesirable bias into the model.

For instance, learning to complete README files may be irrelevant if the model’s primary objective is to excel solely in code completion. However, incorporating various file formats into the context is justified if they hold relevance to the completion task (e.g., including documentation).

This issue can be mitigated through the application of gradient masking. By setting the individual losses of non-completion tokens to zero, these tokens can be excluded from the gradient computation during training.

## 3.3 Fill-in-the-Middle

Due to the autoregressive nature of decoder-only Transformers, they are unable to utilize future tokens in their context. Consequently, to account for both the prefix and suffix of the completion file, the fill-in-the-middle approach was proposed by Bavarian et al. (2022), with its origins in the works of Donahue et al. (2020), Aghajanyan et al. (2022), and Fried et al. (2022). This capability is particularly advantageous for the completion task, as code is frequently written in a non-sequential and chaotic order.

The primary concept of FIM involves randomly dividing a portion of training sequences into three segments: prefix, middle, and suffix. These segments are then concatenated using special tokens to form a new sequence order: prefix, suffix, and middle. Incorporating such augmented data into the pre-training process, introduces a new capability to the model, albeit with a marginal performance degradation (Allal et al. 2023; Rozière et al. 2023).

## 3.4 Evaluation

### 3.4.1 Metrics

### Exact Match

The exact match (EM) metric is a fundamental and widely used measure for evaluating code completion. It is defined as the ratio of correctly completed code lines to the total number of lines. This metric is particularly valued for its direct alignment with the objectives of code completion evaluation.

However, the exact match metric operates at a line-level granularity, which means it only indicates whether a line is completed correctly or not. This binary nature of the metric does not offer insights into the degree of deviation from the correct completion in cases where the completion is incorrect. To mitigate this limitation, the exact match metric is often used in conjunction with the more granular measures.

### Edit Similarity

The generation of functionally correct code completions is valuable, yet the effort required to edit and adapt these generated lines is equally significant. Edit similarity (ES) quantifies this effort by measuring the number of single-character edits (insertions, deletions, or substitutions) needed to transform the generated code into a reference (Svyatkovskiy et al. 2020). Mathematically, ES for two strings is expressed as:

$$\text{ES}(\mathbf{x}, \mathbf{y}) = 1 - \frac{\text{lev}(\mathbf{x}, \mathbf{y})}{\max\{|\mathbf{x}|, |\mathbf{y}|\}}, \quad (3.1)$$

where  $\text{lev}(\mathbf{x}, \mathbf{y})$  represents the Levenshtein distance between the generated code  $\mathbf{x}$  and the reference sample  $\mathbf{y}$ , and  $|\cdot|$  denotes the length of the string.

This metric is crucial for evaluating code completion scenarios, as it provides a measure of the effort developers must exert to correct errors in the generated code. Moreover, it has been shown that edit similarity moderately correlates with the generation of functionally correct code (Dibia et al. 2022).

### Cross-Entropy and Perplexity

As mentioned earlier, cross-entropy is employed as a loss function in the training process of language models. More specifically, it represents the average log-likelihood of the individual ground truth tokens. It can also be interpreted as the average number of nats required to encode the model’s predictions per token compared to the true distribution.

Perplexity (PPL) is the exponentiated form of cross-entropy, providing a more intuitive interpretation as the average number of choices among which the model is uncertain. For this reason, perplexity is often referred to as the weighted average branching factor of a language (Murphy 2022).

Both metrics are frequently used as proxies for assessing model quality. They are valuable for monitoring because these measures are consistently computed during training, being either derived from the loss function (PPL) or

representing the loss itself (CE). However, they are too abstract to serve as primary metrics for specific tasks. Additionally, these measures are heavily influenced by vocabulary size and tokenization methods, which makes them unsuitable for comparing different models.

### Top- $k$ Accuracy

Top- $k$  accuracy is a metric that quantifies the frequency with which the model’s top- $k$  predictions align with the actual ground truth completion. Within the scope of this thesis, top- $k$  accuracy is considered as a token-level metric to provide a more granular evaluation of the model’s performance.

### Pass@ $k$

All aforementioned metrics are syntax-based, meaning they evaluate the model’s performance based on the syntactic match between the generated and reference completions. However, in real-world scenarios, there exists a wide variety of correct completions that may not be present in the dataset used for evaluation. The model might predict a completion that fulfills the user’s functional requirements, even if the metric evaluates it as incorrect.

To address this issue, an unbiased estimation of the probability that at least one generated solution out of  $k$  passes all unit tests was proposed by Chen et al. (2021). This metric, known as pass@ $k$ , is calculated using the following expectation:

$$\text{pass@}k = \mathbb{E}_{\text{Problems}} \left[ 1 - \frac{\binom{n-c}{k}}{\binom{n}{k}} \right], \quad (3.2)$$

where  $n \geq k$  is the number of samples generated by the model, and  $c \leq n$  is the number of those samples that passed all unit tests. This metric reflects the probability of finding a correct solution within  $k$  attempts, which mirrors the iterative approach developers often take when exploring multiple solutions.

To apply this metric to single-line code completion, it is assumed that all other lines, except the target one, are present and functionally correct. This assumption does not hold for many use cases, such as when a user writes a function on the fly. Therefore, pass@ $k$  offers an optimistic assessment of completion capabilities, as the evaluation set contains one corrupted line per problem. In addition to this, Liu et al. (2024) introduced a comprehensive list of further limitations in their Appendix A.4 section. In summary, pass@ $k$  is more realistic for tasks with greater completion granularity or for code generation assessment.

### Further Metrics

Several metrics for evaluating generated code have been migrated from machine translation and summarization fields, yet they often face limitations

due to their unnatural application to code (Evtikhiev et al. 2022). This has prompted the research community to develop code-specific adaptations. However, identifying an optimal metric for code-oriented tasks remains an open challenge.

***N*-gram Based Metrics:** The bilingual evaluation understudy (BLEU) (Papineni et al. 2002) is a precision-oriented metric that measures the proportion of *n*-grams in the candidate text that also appear in the reference snippet. BLEU4 is a specific instance of BLEU that considers *n*-grams up to a length of 4. This configuration is widely adopted because it balances the evaluation of both individual words and longer phrases. The recall-oriented understudy for gisting evaluation (ROUGE) (Lin 2004) is a family of recall-based metrics originally developed for evaluating automatic summarization. ROUGE-N calculates *n*-gram recall between a candidate and a reference, while ROUGE-L uses the longest common subsequence to capture word order without requiring consecutive matches.

**Character-Level Metric:** The ChrF metric (Popović 2015) is an *F*-measure that evaluates character *n*-grams instead of word ones. It incorporates a  $\beta$  parameter to adjust the emphasis between precision and recall. ChrF represents an advancement over previously mentioned metrics by balancing precision and recall in a single measure. The character-level approach offers several advantages: it eliminates tokenization dependencies, implicitly captures morpho-syntactic information, and requires no additional linguistic tools.

**Semantic-Enhanced Metrics:** Metric for evaluation of translation with explicit ordering (METEOR) (Banerjee et al. 2005) improves upon simple *n*-gram matching by adding support for stemming, synonymy, and paraphrasing, along with a penalty for disordered *n*-grams. BERTScore (Zhang et al. 2019) leverages contextual embeddings from BERT (Bidirectional Encoder Representations from Transformers) to compute similarity between tokens in candidate and reference texts, enabling robust matching of paraphrases and semantic equivalents.

**Code-Specific Metrics:** RUBY (Tran et al. 2019) was specifically designed for code evaluation, integrating lexical, syntactic, and semantic representations of source code to better capture functional equivalence. It uses an ensemble approach that leverages the most reliable available representation: program dependence graphs, abstract syntax trees, and surface text. CodeBLEU (Ren et al. 2020) extends BLEU by adding weighted *n*-gram matching for programming language keywords, an abstract syntax tree match component to capture structural information, and a data-flow match to evaluate semantic correctness.

### 3.4.2 Evaluation Modes

Besides the chosen inference strategy, types of which were discussed in Section 2.5, there are multiple setups for evaluating the repository-level completion capabilities of the Code LLM.

During training, language models typically employ a technique known as teacher forcing. In this method, each subsequent token prediction is conditioned on the ground truth tokens from the preceding sequence. This approach takes advantage of the parallelizable nature of Transformers, facilitating efficient training. Conversely, during evaluation, the model can be conditioned on its own previously generated tokens, offering a more realistic assessment of its capabilities. Despite this, teacher forcing remains prevalent during training for the purpose of collecting metrics and monitoring progress, as it does not require additional forward passes. It is crucial to recognize that teacher forcing can render some metrics meaningless, as they may assume conditioning on the model’s output distribution. For example, edit similarity necessitates the generation of a full line solely by the model without external guidance, as it compares entire lines rather than individual tokens. The same statement does not hold for exact match, where a single incorrectly generated token results in the entire line being deemed incorrect.

For single-line code completion, the repository context can be composed based on the file prefix (and suffix in the case of FIM) or the entire file. The former method requires a separate composition and forward pass for each target line in the composition file, offering superior performance. In contrast, the latter requires only a single instance of both, but it risks data leakage of the target line for the retrieval mechanism. Despite the potential issues of the second method, it is suitable for training as it does not necessitate gradient masking for the file prefix and provides more training tokens for loss computation.

Apart from training evaluation, a validation loop is employed periodically to monitor the decline in a model’s ability to generalize to new data, a phenomenon known as overfitting, which occurs due to excessive memorization of the training data. To ensure that validation metrics are comparable to training ones, the same evaluation setup is often utilized. However, this approach differs from the final evaluation of the model on various benchmarks, which is conducted using scenarios that closely resemble real-world conditions.

### 3.4.3 Benchmarks

In recent years, the field of code completion has witnessed significant advancements in benchmarking, particularly for repository-level tasks. This section presents a chronological overview of major benchmarks that have shaped this domain.

CodeXGLUE, introduced by Lu et al. (2021), represents one of the first

comprehensive benchmarks for code intelligence. It encompasses a collection of 10 tasks across 14 datasets, including code completion, which it evaluates using datasets like PY150 (Raychev et al. 2016) and GitHub Java Corpus (Allamanis et al. 2013). While CodeXGLUE established a foundation for code completion evaluation, it primarily focused on in-file completion without considering cross-file dependencies that are prevalent in real-world software development.

Addressing this limitation, Ding et al. (2022) proposed CoCoMIC, which pioneered the joint modeling of in-file and cross-file context for code completion. CoCoMIC’s key innovation was the introduction of a cross-file context finder (CCFINDER) that locates and retrieves relevant cross-file context. This approach significantly improved completion accuracy, particularly for API usage scenarios, highlighting the importance of repository-level understanding in code completion tasks.

Building upon these cross-file context concepts, Zhang et al. (2023) introduced RepoEval, a benchmark constructed from high-quality real-world repositories. Unlike its predecessors, RepoEval explicitly focused on three levels of code completion granularity: line, API invocation, and function body completion scenarios. The benchmark also employs unit tests present in the repositories for evaluating function body completions. RepoEval’s design brought the evaluation closer to real-world development practices.

RepoBench, proposed by Liu et al. (2023), address RepoEval’s limitation of limited repository coverage by significantly expanding the evaluation dataset. It further enhanced the evaluation framework by decomposing the repository-level code completion process into three interconnected tasks: RepoBench-R (retrieval), RepoBench-C (code completion with both 2K and 8K context lengths), and RepoBench-P (end-to-end pipeline). This modular approach enabled more targeted evaluations of each component while maintaining their interdependence in a complete system. RepoBench also expanded language coverage to include both Python and Java, representing a step toward multilingual evaluation. In addition, RepoBench demonstrated that incorporating cross-file contexts significantly improves code completion performance, even with randomly selected contexts.

CrossCodeEval (Ding et al. 2023) significantly broadened the multilingual scope by incorporating four popular programming languages: Python, Java, TypeScript, and C#. It introduced a rigorous methodology for creating examples that strictly require cross-file context for accurate completion, utilizing static analysis to pinpoint cross-file context usage within the current file. Their evaluations demonstrated that models performed poorly with only in-file context but improved significantly when cross-file context was included, suggesting BM25 (Robertson et al. 2009) as an effective retrieval method. This benchmark established a higher standard for evaluating cross-file contextual understanding in diverse programming languages. Later, CrossCodeLongEval (Wu et al. 2024a) expanded upon this foundation by addressing CrossCodeEval’s limited task coverage, introducing chunk and function completion sce-

narios alongside line completion to provide a more comprehensive evaluation framework.

Advancing toward more realistic scenarios, Deng et al. (2024) proposed  $R^2C^2$ -Bench, which introduced a context perturbation strategy to simulate real-world repository-level code completion environments. This benchmark constructs candidate retrieval pools with both abstract and snippet contexts, capturing both coarse-grained global information and fine-grained local details.  $R^2C^2$ -Bench’s comprehensive approach highlighted limitations in previous benchmarks, particularly regarding the coverage of diverse completion scenarios.

Long Code Arena (LCA), as introduced by Bogomolov et al. (2024), presented a comprehensive suite of six benchmarks, with a particular focus on evaluating repository-level capabilities across multiple code-related tasks, including project-level code completion. This component of LCA offers a more nuanced evaluation of completion tasks by categorizing target lines into six distinct types. Notably, the two most significant categories are *infile* and *in-project*, which respectively denote lines that utilize an API declared within the completion file and in other files within the repository. The dataset is stratified into four segments based on the total number of characters forming the overall context, thereby representing varying levels of complexity. To prevent data leakage, a traversal of Git history was conducted. In summary, Long Code Arena added an additional layer of diversity to repository-level code completion benchmarks, enhancing the evaluation landscape.

Wu et al. (2024b) proposed RepoMasterEval, which prioritized testing in authentic development conditions. Recognizing limitations in test suite quality of previous benchmarks, RepoMasterEval employed mutation testing and manual test case crafting to ensure robust evaluation. Uniquely, this benchmark conducted an industrial study correlating model performance with online acceptance rates, demonstrating that RepoMasterEval scores can indeed predict real-world usability of code completion systems.

Further emphasizing developer experience, Codev-Bench (Pan et al. 2024) introduced a developer-centric evaluation framework based on business analysis from industrial code completion products. It redefined evaluation criteria to better align with developers’ intent and desired completion behavior throughout the coding process. Codev-Bench leveraged an agent-based system to automate repository crawling, environment construction, and dynamic call chain extraction from existing unit tests, providing a more realistic assessment of code completion in modern software development.

Most recently,  $M^2RC$ -EVAL (Liu et al. 2024) significantly expanded the multilingual coverage to 18 programming languages, addressing the limited language scope of previous benchmarks. It introduced two types of fine-grained annotations (bucket-level and semantic-level) that enable comprehensive analysis of model performance across different code semantics and complexity levels.  $M^2RC$ -EVAL also created a massively multilingual instruction corpus to



enhance repository-level code completion capabilities of existing models, establishing a new standard for comprehensive multilingual evaluation.

The evolution of these benchmarks reflects the field’s progressive understanding of the complexities inherent in repository-level code completion, moving from isolated single-file evaluations to comprehensive assessments that consider cross-file dependencies, multilingual capabilities, and real-world usage patterns. These benchmarks collectively provide a robust framework for evaluating and improving code completion systems in increasingly realistic scenarios.

### **3.5** Modern Code Completion Models

There will be an overview of the recent Code LLMs that are related to the thesis.

## In-Context Learning

There will be:

- Definition of in-context learning (ICL),
- References to studies supporting the usefulness of ICL for repository-level code completion capabilities of Code LLMs,
- Discussion on prompt engineering,
- Comparison of zero-shot, one-shot, and few-shot learning.

## Chapter 5

# Long Context

The following topics will be covered:

- The significance of long context for in-context learning,
- An overview of the challenges associated with long context, such as the quadratic dependency of computational resources on context length, the inability of RoPE to extrapolate, and the noise generated by softmax-based attention,
- Strategies for addressing long context challenges, including memory-efficient attention implementations, key-value (KV) caching, gradient checkpointing, retrieval-augmented generation (RAG), subquadratic attention, and context extension methods.

**Part II**

**Applied Research**

# Technical Foundation

This chapter offers a comprehensive description of the context composition framework and the fine-tuning pipeline employed in the research presented in this thesis. Additionally, it includes a list of tools necessary for executing their implementations.

## 6.1 Context Composition Framework

This section begins with an explanation of the independent building blocks of the context composition framework and their integration to form a complete compositional strategy. It then enumerates all the composers utilized in the experimental section, accompanied by their descriptions. Finally, it offers an overview of the additional parameter applied to modify the retrieval order and includes a commentary that invites further exploratory investigation into the framework.


### 6.1.1 Building Blocks

As stated in Section 3.1, we define the context composer as a function that takes a repository snapshot and a completion file and returns a string that represents the project context. This function is decomposed into seven abstract blocks:

1. **File Filtering** determines which files to include in the subsequent data processing pipeline. For instance, it can filter out files with empty content or files with certain extensions.
2. **File Preprocessing** modifies the content of the files. For example, it can extract code from markdown files.

3. **File Chunking** changes the granularity of the data from files to text chunks. The simplest form of chunking is identity, which maintains the file-grained structure.
4. **Chunk Ranking** determines the relevance of the chunks for the completion file. It can take multiple forms, ranging from simple heuristics and sparse embedding comparisons to more sophisticated dense approaches.
5. **Chunk Sorting** defines the rules for ordering the chunks based on their ranks. The most common scenario is to sort the chunks in ascending order of their ranks, placing the most relevant chunks at the end of the list.
6. **Chunk Assembling** joins sorted chunks into a single string using a specific template. For example, the assembler can prepend a comment with path information of the chunk's source file and join them using a separator string.
7. **Context Postprocessing** modifies the context as a unified string. For instance, line dropout can be applied by this block.

Each of these blocks has its own abstract class and serves as a foundation for concrete implementations. This design allows for several different transformations associated with a single block type to be part of a single context composer, eliminates code duplication and allows for greater flexibility in composer creation. The directed graph presented in Figure 6.1 visualizes the order in which these instances can be applied.



■ **Figure 6.1** Order of composer blocks. Nodes represent block types: ① File Filtering, ② File Preprocessing, ③ File Chunking, ④ Chunk Ranking, ⑤ Chunk Sorting, ⑥ Chunk Assembling, and ⑦ Context Postprocessing. Edges indicate the permissible sequence of block application. The data format is denoted by three frames: Files, Chunks, and Context. The bold subgraph highlights blocks that may be omitted.

### 6.1.2 List of Context Composers

All context composers utilized in the Research Investigation chapter adhere to two standard preprocessing steps: the removal of empty files and the nor-

malization of all line endings to the Line Feed (LF) character. These shared preprocessing steps ensure consistency across all experiments. The following list provides a comprehensive overview of the context composers employed in this research.


1. **File-Level** produces an empty context.
2. **Path Distance .py** constructs the context using only files with the `.py` extension. The selected files are sorted in descending order according to their path distance from the completion file. For files with identical path distances, a secondary sorting criterion is applied using the Intersection over Union (IoU) metric, which equals to the number of lines shared with the completion file divided by the total number of unique lines in both files. The IoU calculation considers lines with leading and trailing whitespace removed and includes only those lines that are at least five characters in length after whitespace removal.
3. **Lines IoU .py** is similar to the Path Distance .py method, but omits the initial sorting by path distance. Instead, files are ranked directly using the IoU metric.
4. **Code Chunks** removes all docstrings, comments, and import statements from the context produced by Path Distance .py.
5. **Half-Memory .py** begins with the context generated by Path Distance .py. Each line is independently removed with a dropout probability of 0.5, maintaining the overall saturation of the context window.
6. **Declarations .py** extends Path Distance .py by filtering out all non-declarative elements, retaining only function and class declarations.
7. **Text Chunks .py** uses Path Distance .py as the base composer, but removes all code from the context, leaving only docstrings and comments.
8. **Text Files** constructs the context using files with the extensions `.json`, `.yaml`, `.yml`, `.sh`, `.md`, `.txt`, and `.rst`. The selected files are grouped in ascending order of relevance: `[.json]`, `[.yaml, .yml]`, `[.sh]`, `[.md, .txt, .rst]`. Within each group, files are further sorted in descending order according to their path distance from the completion file.
9. **Random Files** constructs the context by randomly ordering all files from the repository snapshot.
10. **Random .py** selects only files with the `.py` extension and arranges them in a random order.
11. **Random Tokens** constructs the context by sampling a sequence of non-special tokens at random, with each token selected independently and with equal probability.

12. **Duplication** constructs the context by repeatedly concatenating the content of the completion file until the maximum context window size is reached.
13. **Leak** begins with the context produced by Path Distance .py. The completion file is randomly divided into five segments at newline characters, which then disjointedly replace context lines at random positions, approximately preserving the original token count.

### 6.1.3 Order Modifications

To increase the diversity of experiments conducted to address RQ.B1 and RQ.B2, two additional file ordering modifications are considered alongside the *original* mode, as illustrated in Figure 6.2:

- *reversed*: The retrieved files that fit within the model’s context window are reordered in reverse. Consequently, the most relevant file appears at the beginning of the context window.
- *irrelevant*: The order of retrieved files is reversed prior to the context length cut-off, resulting in a complete reversal of the context, with the most irrelevant files positioned at the end of the context string.



■ **Figure 6.2** Three modes of context ordering. The numbers in boxes indicate the rank of the chunks, with smaller numbers denoting higher relevance.

### 6.1.4 Examples

In addition to the blocks used to create composers employed in our experiments, we provide further exemplary implementations in the attached code, totaling 30 distinct blocks. The repository also includes a demonstrative Jupyter Notebook that describes these blocks and offers concrete usage examples of the package.



## 6.2 Fine-Tuning Pipeline

For the purposes of this research, we developed a fine-tuning pipeline aimed at adapting Code LLMs to the repository-level context.

From the data perspective, it supports the integration of pre-composed datasets obtained from the aforementioned context composition framework. The train-validation split is designed to eliminate data leakage from samples gathered from the same repositories. Configurable tokenization ensures context window saturation and facilitates the application of gradient masking.

The pipeline integrates models from the Hugging Face Hub<sup>1</sup>. It supports adapters that can augment the model architecture, freeze a subset of parameters, or modify the optimizer initialization. Models are dynamically loaded onto a single device based on hardware availability. Efficient attention implementation and parameter data types are selected using the same principle.

The training loop is implemented in pure PyTorch<sup>2</sup>. It supports cosine learning rate scheduling with linear warmup, gradient scaling, clipping, and accumulation. The pipeline can periodically invoke two validation loops: one on the original dataset split and another on an additional provided dataset. The latter allows metrics to be grounded on baseline composition strategy instead of solely monitoring the in-distribution capabilities of the model.

During the training process, the pipeline produces a set of configurable metrics and statistics for both the training and validation data. These include cross-entropy, exact match, top- $k$  accuracy, learning rate tracking, epoch and token number counting. Their values are decomposed based on the types of tokens they are applied to. We denote five such types:

- **Attached:** Tokens whose losses participate as leaf nodes in automatic differentiation for gradient computation.
- **Detached:** Tokens whose losses are zeroed out during gradient masking and serve as a complement to the attached tokens.
- **Completion:** Tokens of the completion file.
- **Context:** Tokens of the repository context.
- **Full:** All tokens.

In addition to statistics and metrics, the pipeline provides logging of the standard output and error streams. Checkpointing of the model and optimizer states is also supported.

The entire pipeline utilizes a modular structure, which enhances code maintainability and extendability. It is designed in an end-to-end manner, requiring only a pre-composed dataset, a set of YAML configuration files, and a single

---

<sup>1</sup><https://huggingface.co/>

<sup>2</sup><https://pytorch.org/>

command line from the user. All components that employ stochastic behavior are accompanied by a seed parameter, ensuring that all experiments are reproducible.

### 6.3 Tools

In the context composition framework, we integrate Hugging Face tokenizers from the Transformers<sup>3</sup> library to compute token-related statistics, and Tree-sitter<sup>4</sup> to decompose Python code into its parts. Both libraries are employed in the concrete implementations of the blocks, although their usage can be omitted without disrupting the main abstract framework. The OmegaConf<sup>5</sup> configuration system is utilized to manage the initialization of the composers from YAML configuration files.

The pipeline inherits requirements from the context composition framework as it utilizes its functions. Additionally, it employs Datasets<sup>6</sup> and Accelerate<sup>7</sup> from the Hugging Face ecosystem to provide peripheral support to the Transformers library. The Hydra<sup>8</sup> package extends the configurative capabilities of OmegaConf and enhances the pipeline’s usability. PyTorch, as mentioned in the previous section, is used to implement the training loop. The tqdm<sup>9</sup> library is employed to display progress bars for all prolonged processes. Weights & Biases<sup>10</sup> is used to track metrics and statistics in real-time.

To conduct all experiments and run evaluation scripts, a single 8×H200 node was used, with each GPU having 141 GB of VRAM.

---

<sup>3</sup><https://github.com/huggingface/transformers>

<sup>4</sup><https://tree-sitter.github.io/tree-sitter/>

<sup>5</sup><https://github.com/omry/omegaconf>

<sup>6</sup><https://github.com/huggingface/datasets>

<sup>7</sup><https://github.com/huggingface/accelerate>

<sup>8</sup><https://github.com/facebookresearch/hydra>

<sup>9</sup><https://github.com/tqdm/tqdm>

<sup>10</sup><https://github.com/wandb/wandb>

# Research Investigation

In this chapter we describe the research aspect of the thesis.

## 7.1 Research Questions

To investigate the nature of in-context learning capabilities of Code LLMs, we consider the following research questions:

**RQ.A1 Composition Impact on Inference:** Does the quality improvement of code completion depend on the composition strategy employed during model inference?

**RQ.A2 Fine-Tuning on Compositions:** Does fine-tuning a pre-trained Code LLM with a specific context composer enhance the subsequent quality of code completion?

**RQ.B1 Effect of Context Extension:** Does the repository-level pre-training step affect the in-context learning abilities of the model developed during earlier stages?

**RQ.B2 Influence of Composition on Context Extension:** Does the quality improvement of code completion depend on the context composition approach used during the repository-level pre-training step?

## 7.2 Experimental Design

In this section we describe the unified experimental design for all the four research questions while the specific details are deferred to the subsequent sections dedicated to each research question separately.

### 7.2.1 Training Data

The dataset utilized in this research was provided by the supporting company and is based on the methodology outlined in Bogomolov et al. (2024). The dataset was constructed by traversing the GitHub histories of Python repositories and applying permissive license filtering to sample completion files and their corresponding parent commits. Each data point consists of a pair: a list of Python completion files and a repository snapshot that captures the state of the repository at the time the completion files were added. The snapshot includes all text files except for the completion files themselves. To prevent contamination of the benchmark used in the evaluation, any repositories present in the benchmark were excluded from the training data.

The resulting corpus comprises 1,640 repositories, 160,801 commits, and 361,052 completion files. The completion files contain a total of 1.7 billion characters, while the repository snapshot files contain 4.8 trillion characters.

To delineate the scope of this work, it is important to note that the corpus described thus far was entirely provided by the company and was not generated by the author of this thesis. Conversely, the subsequent processing steps applied to this dataset represent the original contributions of the author.

### 7.2.2 Training

### 7.2.3 Evaluation

## 7.3 Composition Impact on Inference

### 7.4 Fine-Tuning on Compositions

### 7.5 Effect of Context Extension

### 7.6 Influence of Composition on Context Extension

### 7.7 Limitations and Future Work

# Conclusion

In conclusion, the objectives of this thesis have been accomplished. The theory of modern repository-level code completion systems is covered in Conceptual Framework, providing an overview of the code completion task, a general formulation of language modeling, and deep learning methods addressing this area. Subsequently, the specifics of project-level completion are discussed through the lens of language modeling. The relevance and importance of the in-context learning phenomenon are described. To further explore the nuances of repository-level code completion, the general problems of long context are presented along with methods to address them. The research aspect of this work is detailed in Applied Research, which includes answers to four research questions and a description of the practical framework developed to support these investigations.

The conducted research has yielded several significant findings regarding repository-level code completion. The results substantiate the crucial role of repository context in enhancing completion quality. Contemporary base Code LLMs possess sufficient training, with their in-context learning capabilities showing only marginal improvement through downstream fine-tuning on different context composition strategies. Data leakage in the repository context during the context extension phase adversely affects the model’s in-context abilities, while other composers maintain baseline performance integrity. The context composition strategy employed during the repository-level pre-training stage demonstrates minimal impact on final model quality, suggesting that RoPE adjustment serves as the primary driver of long-context improvements and underscores its usage limitations. Additionally, computational requirements of the repository-level pre-training stage can be substantially reduced while maintaining competitive results. For instance, file-level training, even without repository context, remains highly effective. Finally, gradient masking produces a statistically significant yet marginal performance decrease on composers generating inlier context.

Overall, this thesis contributes to the evolution of code completion models by underscoring theoretical foundations in the conceptual framework and

offering valuable empirical insights for practitioners through comprehensive research on context composition techniques for repository-level understanding.



## Appendix A

# Appendix

Not ready.

- Aghajanyan, Armen et al. (2022). *CM3: A Causal Masked Multimodal Model of the Internet*. arXiv: 2201.07520 [cs.CL]. URL: <https://arxiv.org/abs/2201.07520>.
- Allal, Loubna Ben et al. (2023). *SantaCoder: don't reach for the stars!* arXiv: 2301.03988 [cs.SE]. URL: <https://arxiv.org/abs/2301.03988>.
- Allamanis, Miltiadis and Charles Sutton (2013). “Mining source code repositories at massive scale using language modeling”. In: *2013 10th Working Conference on Mining Software Repositories (MSR)*, pp. 207–216. DOI: 10.1109/MSR.2013.6624029.
- Alon, Uri et al. (2019). *Structural Language Models of Code*. arXiv: 1910.00577 [cs.LG]. URL: <https://arxiv.org/abs/1910.00577>.
- Asaduzzaman, Muhammad et al. (2014). “Context-Sensitive Code Completion Tool for Better API Usability”. In: *2014 IEEE International Conference on Software Maintenance and Evolution*, pp. 621–624. DOI: 10.1109/ICSME.2014.110.
- Ba, Jimmy Lei, Jamie Ryan Kiros, and Geoffrey E. Hinton (2016). *Layer Normalization*. arXiv: 1607.06450 [stat.ML]. URL: <https://arxiv.org/abs/1607.06450>.
- Baevski, Alexei and Michael Auli (2018). *Adaptive Input Representations for Neural Language Modeling*. arXiv: 1809.10853 [cs.CL]. URL: <https://arxiv.org/abs/1809.10853>.
- Bahdanau, Dzmitry, Kyunghyun Cho, and Yoshua Bengio (2014). *Neural Machine Translation by Jointly Learning to Align and Translate*. arXiv: 1409.0473 [cs.CL]. URL: <https://arxiv.org/abs/1409.0473>.
- Bakal, Gal et al. (2025). *Experience with GitHub Copilot for Developer Productivity at Zoominfo*. arXiv: 2501.13282 [cs.SE]. URL: <https://arxiv.org/abs/2501.13282>.
- Banerjee, Satanjeev and Alon Lavie (June 2005). “METEOR: An Automatic Metric for MT Evaluation with Improved Correlation with Human Judgments”. In: *Proceedings of the ACL Workshop on Intrinsic and Extrinsic Evaluation Measures for Machine Translation and/or Summarization*. Ann Arbor, Michigan: Association for Computational Linguistics, pp. 65–72. URL: <https://www.aclweb.org/anthology/W05-0909>.
- Barbero, Federico et al. (2024). *Round and Round We Go! What makes Rotary Positional Encodings useful?* arXiv: 2410.06205 [cs.CL]. URL: <https://arxiv.org/abs/2410.06205>.
- Bavarian, Mohammad et al. (2022). *Efficient Training of Language Models to Fill in the Middle*. arXiv: 2207.14255 [cs.CL]. URL: <https://arxiv.org/abs/2207.14255>.
- Bielik, Pavol, Veselin Raychev, and Martin Vechev (20–22 Jun 2016). “PHOG: Probabilistic Model for Code”. In: *Proceedings of The 33rd International Conference on Machine Learning*. Ed. by Maria Florina Balcan and Kilian Q. Weinberger. Vol. 48. Proceedings of Machine Learning Research. New



- York, New York, USA: PMLR, pp. 2933–2942. URL: <https://proceedings.mlr.press/v48/bielik16.html>.
- Bogomolov, Egor et al. (2024). *Long Code Arena: a Set of Benchmarks for Long-Context Code Models*. arXiv: 2406.11612 [cs.LG]. URL: <https://arxiv.org/abs/2406.11612>.
- Chen, Mark et al. (2021). *Evaluating Large Language Models Trained on Code*. arXiv: 2107.03374 [cs.LG]. URL: <https://arxiv.org/abs/2107.03374>.
- Child, Rewon et al. (2019). *Generating Long Sequences with Sparse Transformers*. arXiv: 1904.10509 [cs.LG]. URL: <https://arxiv.org/abs/1904.10509>.
- Ciniselli, Matteo et al. (2021). “An Empirical Study on the Usage of Transformer Models for Code Completion”. In: *IEEE Transactions on Software Engineering*, pp. 1–1. ISSN: 2326-3881. DOI: 10.1109/tse.2021.3128234. URL: <http://dx.doi.org/10.1109/TSE.2021.3128234>.
- Deng, Ken et al. (2024). *R<sup>2</sup>C<sup>2</sup>-Coder: Enhancing and Benchmarking Real-world Repository-level Code Completion Abilities of Code Large Language Models*. arXiv: 2406.01359 [cs.CL]. URL: <https://arxiv.org/abs/2406.01359>.
- Dibia, Victor et al. (2022). *Aligning Offline Metrics and Human Judgments of Value for Code Generation Models*. arXiv: 2210.16494 [cs.SE]. URL: <https://arxiv.org/abs/2210.16494>.
- Ding, Yangruibo et al. (2022). *CoCoMIC: Code Completion By Jointly Modeling In-file and Cross-file Context*. arXiv: 2212.10007 [cs.CL]. URL: <https://arxiv.org/abs/2212.10007>.
- Ding, Yangruibo et al. (2023). *CrossCodeEval: A Diverse and Multilingual Benchmark for Cross-File Code Completion*. arXiv: 2310.11248 [cs.LG]. URL: <https://arxiv.org/abs/2310.11248>.
- Donahue, Chris, Mina Lee, and Percy Liang (2020). *Enabling Language Models to Fill in the Blanks*. arXiv: 2005.05339 [cs.CL]. URL: <https://arxiv.org/abs/2005.05339>.
- Estoup, J. B. (1916). “Les Gammes Stenographiques”. In: *Institut Stenographique de France*.
- Evtikhiev, Mikhail et al. (Sept. 2022). “Out of the BLEU: How should we assess quality of the Code Generation models?” In: *Journal of Systems and Software* 203, p. 111741. ISSN: 0164-1212. DOI: 10.1016/j.jss.2023.111741. URL: <http://dx.doi.org/10.1016/j.jss.2023.111741>.
- Fried, Daniel et al. (2022). *InCoder: A Generative Model for Code Infilling and Synthesis*. arXiv: 2204.05999 [cs.SE]. URL: <https://arxiv.org/abs/2204.05999>.
- Gage, Philip (1994). “A new algorithm for data compression”. In: *The C Users Journal archive* 12, pp. 23–38. URL: <https://api.semanticscholar.org/CorpusID:59804030>.

- Giagnorio, Alessandro, Alberto Martin-Lopez, and Gabriele Bavota (2025). *Why Personalizing Deep Learning-Based Code Completion Tools Matters*. arXiv: 2503.14201 [cs.SE]. URL: <https://arxiv.org/abs/2503.14201>.
- Ginzberg, Adam, Lindsey Kostas, and Tara Balakrishnan (2017). “Automatic Code Completion”. In: URL: <https://api.semanticscholar.org/CorpusID:35359813>.
- Guo, Daya et al. (2024). *DeepSeek-Coder: When the Large Language Model Meets Programming – The Rise of Code Intelligence*. arXiv: 2401.14196 [cs.SE]. URL: <https://arxiv.org/abs/2401.14196>.
- Han, Sangmok, David R. Wallace, and Robert C. Miller (2009). “Code Completion from Abbreviated Input”. In: *2009 IEEE/ACM International Conference on Automated Software Engineering*, pp. 332–343. DOI: 10.1109/ASE.2009.64.
- Harris, Zellig S. (1954). “Distributional Structure”. In: URL: <https://api.semanticscholar.org/CorpusID:86680084>.
- Hendrycks, Dan and Kevin Gimpel (2016). *Gaussian Error Linear Units (GELUs)*. arXiv: 1606.08415 [cs.LG]. URL: <https://arxiv.org/abs/1606.08415>.
- Hill, Rosco and Joe Rideout (2004). “Automatic Method Completion”. In: *Proceedings of the 19th IEEE International Conference on Automated Software Engineering*. ASE ’04. USA: IEEE Computer Society, pp. 228–235. ISBN: 0769521312.
- Hindle, Abram et al. (2012). “On the naturalness of software”. In: *Proceedings of the 34th International Conference on Software Engineering*. ICSE ’12. Zurich, Switzerland: IEEE Press, pp. 837–847. ISBN: 9781467310673.
- Holtzman, Ari et al. (2019). *The Curious Case of Neural Text Degeneration*. arXiv: 1904.09751 [cs.CL]. URL: <https://arxiv.org/abs/1904.09751>.
- Huang, Siming et al. (2024). *OpenCoder: The Open Cookbook for Top-Tier Code Large Language Models*. arXiv: 2411.04905 [cs.CL]. URL: <https://arxiv.org/abs/2411.04905>.
- Hui, Binyuan et al. (2024). *Qwen2.5-Coder Technical Report*. arXiv: 2409.12186 [cs.CL]. URL: <https://arxiv.org/abs/2409.12186>.
- Karampatsis, Rafael-Michael and Charles Sutton (2019). *Maybe Deep Neural Networks are the Best Choice for Modeling Source Code*. arXiv: 1903.05734 [cs.SE]. URL: <https://arxiv.org/abs/1903.05734>.
- Kim, Seohyun et al. (2021). *Code Prediction by Feeding Trees to Transformers*. arXiv: 2003.13848 [cs.SE]. URL: <https://arxiv.org/abs/2003.13848>.
- Kingma, Diederik P. and Jimmy Ba (2014). *Adam: A Method for Stochastic Optimization*. arXiv: 1412.6980 [cs.LG]. URL: <https://arxiv.org/abs/1412.6980>.
- Lin, Chin-Yew (July 2004). “ROUGE: A Package for Automatic Evaluation of Summaries”. In: *Text Summarization Branches Out*. Barcelona, Spain: Association for Computational Linguistics, pp. 74–81. URL: <https://aclanthology.org/W04-1013/>.

- Liu, Fang et al. (2020). *Multi-task Learning based Pre-trained Language Model for Code Completion*. arXiv: 2012.14631 [cs.SE]. URL: <https://arxiv.org/abs/2012.14631>.
- Liu, Jiaheng et al. (2024). *M<sup>2</sup>RC-EVAL: Massively Multilingual Repository-level Code Completion Evaluation*. arXiv: 2410.21157 [cs.CL]. URL: <https://arxiv.org/abs/2410.21157>.
- Liu, Tianyang, Canwen Xu, and Julian McAuley (2023). *RepoBench: Benchmarking Repository-Level Code Auto-Completion Systems*. arXiv: 2306.03091 [cs.CL]. URL: <https://arxiv.org/abs/2306.03091>.
- Loshchilov, Ilya and Frank Hutter (2017). *Decoupled Weight Decay Regularization*. arXiv: 1711.05101 [cs.LG]. URL: <https://arxiv.org/abs/1711.05101>.
- Lu, Shuai et al. (2021). *CodeXGLUE: A Machine Learning Benchmark Dataset for Code Understanding and Generation*. arXiv: 2102.04664 [cs.SE]. URL: <https://arxiv.org/abs/2102.04664>.
- Mandelin, David et al. (2005). “Jungloid mining: helping to navigate the API jungle”. In: *Proceedings of the 2005 ACM SIGPLAN Conference on Programming Language Design and Implementation*. PLDI ’05. Chicago, IL, USA: Association for Computing Machinery, pp. 48–61. ISBN: 1595930566. DOI: 10.1145/1065010.1065018. URL: <https://doi.org/10.1145/1065010.1065018>.
- Murphy, Kevin P. (2022). *Probabilistic Machine Learning: An introduction*. MIT Press. URL: <http://probml.github.io/book1>.
- Pan, Zhenyu et al. (2024). *Codev-Bench: How Do LLMs Understand Developer-Centric Code Completion?* arXiv: 2410.01353 [cs.SE]. URL: <https://arxiv.org/abs/2410.01353>.
- Papineni, Kishore et al. (2002). “BLEU: a method for automatic evaluation of machine translation”. In: *Proceedings of the 40th Annual Meeting on Association for Computational Linguistics*. ACL ’02. Philadelphia, Pennsylvania: Association for Computational Linguistics, pp. 311–318. DOI: 10.3115/1073083.1073135. URL: <https://doi.org/10.3115/1073083.1073135>.
- Peng, Sida et al. (2023). *The Impact of AI on Developer Productivity: Evidence from GitHub Copilot*. arXiv: 2302.06590 [cs.SE]. URL: <https://arxiv.org/abs/2302.06590>.
- Popović, Maja (Sept. 2015). “chrF: character n-gram F-score for automatic MT evaluation”. In: *Proceedings of the Tenth Workshop on Statistical Machine Translation*. Ed. by Ondřej Bojar et al. Lisbon, Portugal: Association for Computational Linguistics, pp. 392–395. DOI: 10.18653/v1/W15-3049. URL: <https://aclanthology.org/W15-3049/>.
- Proksch, Sebastian, Johannes Lerch, and Mira Mezini (Dec. 2015). “Intelligent Code Completion with Bayesian Networks”. In: *ACM Trans. Softw. Eng. Methodol.* 25.1. ISSN: 1049-331X. DOI: 10.1145/2744200. URL: <https://doi.org/10.1145/2744200>.

- Raychev et al. (2014). “Code completion with statistical language models”. In: *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation*. PLDI '14. Edinburgh, United Kingdom: Association for Computing Machinery, pp. 419–428. ISBN: 9781450327848. DOI: 10.1145/2594291.2594321. URL: <https://doi.org/10.1145/2594291.2594321>.
- Raychev, Veselin, Pavol Bielik, and Martin Vechev (Oct. 2016). “Probabilistic model for code with decision trees”. In: *SIGPLAN Not.* 51.10, pp. 731–747. ISSN: 0362-1340. DOI: 10.1145/3022671.2984041. URL: <https://doi.org/10.1145/3022671.2984041>.
- Ren, Shuo et al. (2020). *CodeBLEU: a Method for Automatic Evaluation of Code Synthesis*. arXiv: 2009.10297 [cs.SE]. URL: <https://arxiv.org/abs/2009.10297>.
- Robertson, Stephen and Hugo Zaragoza (Apr. 2009). “The Probabilistic Relevance Framework: BM25 and Beyond”. In: *Found. Trends Inf. Retr.* 3.4, pp. 333–389. ISSN: 1554-0669. DOI: 10.1561/15000000019. URL: <https://doi.org/10.1561/15000000019>.
- Rozière, Baptiste et al. (2023). *Code Llama: Open Foundation Models for Code*. arXiv: 2308.12950 [cs.CL]. URL: <https://arxiv.org/abs/2308.12950>.
- Sennrich, Rico, Barry Haddow, and Alexandra Birch (2015). *Neural Machine Translation of Rare Words with Subword Units*. arXiv: 1508.07909 [cs.CL]. URL: <https://arxiv.org/abs/1508.07909>.
- Shazeer, Noam (2020). *GLU Variants Improve Transformer*. arXiv: 2002.05202 [cs.LG]. URL: <https://arxiv.org/abs/2002.05202>.
- Srivastava, Nitish et al. (2014). “Dropout: A Simple Way to Prevent Neural Networks from Overfitting”. In: *Journal of Machine Learning Research* 15.56, pp. 1929–1958. URL: <http://jmlr.org/papers/v15/srivastava14a.html>.
- Su, Jianlin et al. (2021). *RoFormer: Enhanced Transformer with Rotary Position Embedding*. arXiv: 2104.09864 [cs.CL]. URL: <https://arxiv.org/abs/2104.09864>.
- Svyatkovskiy, Alexey et al. (2020). *IntelliCode Compose: Code Generation Using Transformer*. arXiv: 2005.08025 [cs.CL]. URL: <https://arxiv.org/abs/2005.08025>.
- Takerngsaksiri, Wannita et al. (2023). *Students’ Perspective on AI Code Completion: Benefits and Challenges*. arXiv: 2311.00177 [cs.SE]. URL: <https://arxiv.org/abs/2311.00177>.
- Tran, Ngoc et al. (May 2019). “Does BLEU Score Work for Code Migration?” In: *2019 IEEE/ACM 27th International Conference on Program Comprehension (ICPC)*. IEEE, pp. 165–176. DOI: 10.1109/icpc.2019.00034. URL: <http://dx.doi.org/10.1109/ICPC.2019.00034>.
- Vaswani, Ashish et al. (2017). *Attention Is All You Need*. arXiv: 1706.03762 [cs.CL]. URL: <https://arxiv.org/abs/1706.03762>.

- Wang, Qiang et al. (2019). *Learning Deep Transformer Models for Machine Translation*. arXiv: 1906.01787 [cs.CL]. URL: <https://arxiv.org/abs/1906.01787>.
- Weber, Thomas et al. (June 2024). “Significant Productivity Gains through Programming with Large Language Models”. In: *Proc. ACM Hum.-Comput. Interact.* 8.EICS. DOI: 10.1145/3661145. URL: <https://doi.org/10.1145/3661145>.
- Wu, Di et al. (2024a). *Repoformer: Selective Retrieval for Repository-Level Code Completion*. arXiv: 2403.10059 [cs.SE]. URL: <https://arxiv.org/abs/2403.10059>.
- Wu, Qinyun et al. (2024b). *RepoMasterEval: Evaluating Code Completion via Real-World Repositories*. arXiv: 2408.03519 [cs.SE]. URL: <https://arxiv.org/abs/2408.03519>.
- Xiong, Ruibin et al. (2020). *On Layer Normalization in the Transformer Architecture*. arXiv: 2002.04745 [cs.LG]. URL: <https://arxiv.org/abs/2002.04745>.
- Zhang, Biao and Rico Sennrich (2019). *Root Mean Square Layer Normalization*. arXiv: 1910.07467 [cs.LG]. URL: <https://arxiv.org/abs/1910.07467>.
- Zhang, Fengji et al. (2023). *RepoCoder: Repository-Level Code Completion Through Iterative Retrieval and Generation*. arXiv: 2303.12570 [cs.CL]. URL: <https://arxiv.org/abs/2303.12570>.

# Contents of the attachment

Not ready.